

Parallelism (PAR)

Introduction to (shared-memory) parallel architectures

Eduard Ayguadé, Julita Corbalán,
Daniel Jiménez and Gladys Utrera

Computer Architecture Department
Universitat Politècnica de Catalunya

Course 2017/18 (Spring semester)

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures

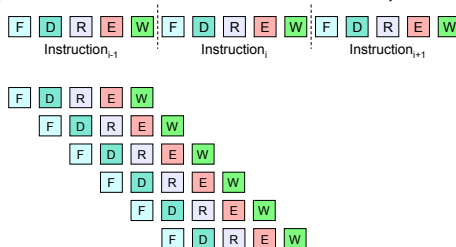
Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Pipelining


- ▶ Execution of single instruction divided in multiple stages
- ▶ Overlap the execution of different stages of consecutive instructions
- ▶ Ideal: $IPC=1$ (1 instruction executed per cycle)




- ▶ $IPC < 1$ due to hazards (structural, data, control), preventing the execution of an instruction in its designated clock cycle

Memory hierarchy

- ▶ Addressing the yearly increasing gap between CPU cycle and memory access times



- ▶ Size vs. access time



- ▶ Non-blocking design

Memory hierarchy

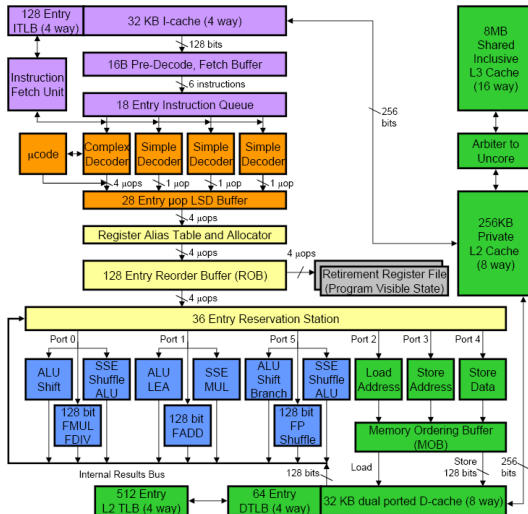
- ▶ The principle of locality: if an item is referenced ...
 - ▶ Temporal locality: ... it will tend to be referenced again soon (e.g., loops, reuse)
 - ▶ Spatial locality: ... items whose addresses are close by tend to be referenced soon (e.g., straight line code, array access)
- ▶ Line (or block)
 - ▶ A number of consecutive words in memory (e.g. 32 bytes, equivalent to 4 words x 8 bytes)
 - ▶ Unit of information that is transferred between two levels in the hierarchy
- ▶ On an access to a level in the hierarchy
 - ▶ Hit: data appears in one of the lines in that level
 - ▶ Miss: data needs to be retrieved from a line in the next level

Sources of parallelism in uniprocessors

- ▶ ILP (Instruction-level parallelism)
 - ▶ Superscalar architecture: multiple issue slots (functional units)
 - ▶ Execution of multiple instructions, from the same instruction flow, per cycle
- ▶ TLP (thread-level parallelism)
 - ▶ Multithreaded architecture¹: fill the pipeline with instructions from multiple instruction flows
 - ▶ Latency hiding (cache misses, non-pipelined FP, ...)
- ▶ DLP (data-level parallelism)
 - ▶ SIMD architecture: single-instruction executed on multiple-data in a single word
 - ▶ Vector functional unit

¹Hyperthreading in Intel terminology

Current uniprocessor architecture: Intel Nehalem i7



Who exploits this uniprocessor parallelism?

In theory, the compiler understands all of this ... but in practice the compiler may need your help:

- ▶ Software pipelining to statically schedule ILP
- ▶ Unrolling to allow the processor to exploit ILP dynamically
- ▶ Data contiguous in memory and aligned to efficiently exploit DLP
- ▶ Blocking (or tiling) to define a problem that fits in register/L1-cache/L2-cache (temporal locality)
- ▶ ...

Reasons and techniques explored in detail in PCA course
(Architecture-Conscious Programming)

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

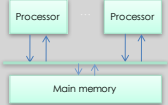
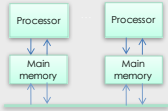
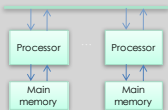
Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms


The memory consistency problem

Classification of multi-processor architectures

Memory architecture	Address space(s)	Connection	Model for data sharing	Names
(Centralized) Shared-memory architecture	Single shared address space, uniform access time		Load/store instructions from processors	<ul style="list-style-type: none"> • SMP (Symmetric Multi-Processor) architecture • UMA (Uniform Memory Access) architecture
Distributed-memory architecture	Single shared address space, non-uniform access time		Load/store instructions from processors	<ul style="list-style-type: none"> • DSM (Distributed-Shared Memory) architecture • NUMA (Non-Uniform Memory Access) architecture
	Multiple separate address spaces		Explicit messages through network interface card	<ul style="list-style-type: none"> • Message-passing multiprocessor • Cluster Architecture • Multicomputer


Symmetric multi-processor architectures

- ▶ Abbreviated SMP
 - ▶ Two or more identical processors are connected to a single shared main memory
 - ▶ Interconnection network: any processor can access to any memory location
- ▶ Symmetric multiprocessing: a single OS instance on the SMP
 - ▶ Asymmetric multiprocessor (e.g. high/low ILP processors, ...) and/or multiprocessing (e.g. some processors running OS, others user code)




Symmetric multi-processor architectures

- ▶ Uniform Memory Access (UMA)
 - ▶ Access to shared data with load/store instructions
 - ▶ Access time to a memory location is independent of which processor makes the request or which memory chip contains the data
- ▶ The bottleneck in the scalability of SMP is the 'bandwidth' of the interconnection network and the memory



Symmetric multi-processor architectures

Local caches and multi-banked (interleaved) memory



The coherence problem




Chart shows value of `foo` (variable stored at address X) stored in main memory and in each processor's cache **

** Assumes write-back cache behavior

Action	P1 \$	P2 \$	P3 \$	P4 \$	mem[X]
					0
P1 load X	0 miss				0
P2 load X	0	0 miss			0
P1 store X	1	0			0
P3 load X	1	0	0 miss		0
P3 store X	1	0	2		0
P2 load X	1	0 hit	2		0
P1 load Y (say this load causes eviction of foo)		0	2		1

(CMU 15-418, Spring 2012)


Coherence protocols

- ▶ Write-update:
 - ▶ Writing processor broadcasts the new value and forces all others to update their copies
 - ▶ Higher bus traffic
- ▶ Write-invalidate:
 - ▶ Writing processor forces all others to invalidate their copies
 - ▶ The new value is provided to others when requested or when flushed from cache

Coherence mechanisms


Snooping:

- ▶ Every cache that has a copy from a block in physical memory keeps its sharing status (**status distributed**)
- ▶ **Broadcast** medium (e.g. a bus) used to make all transactions visible to all caches and define **ordering**
- ▶ Caches monitor or **snoop on the medium** and take action on relevant events (e.g. change status)



Directory-based: the sharing status of each block in memory is kept in just one location (the directory) – to be studied later.

Coherence mechanisms

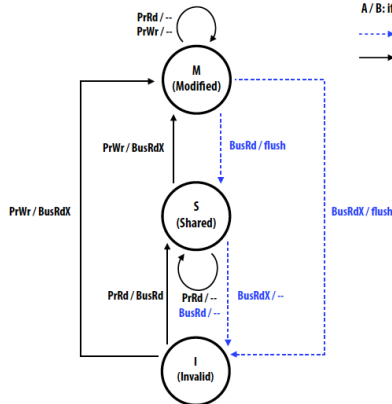


MSI write-invalidate snooping protocol

- ▶ States
 - ▶ Modified (M): only one copy, written (dirty)
 - ▶ Shared (S): one or more copies, all read (clean)
 - ▶ Invalid (I): not valid
- ▶ CPU events
 - ▶ PrRd (Processor read)
 - ▶ PrWr (Processor write)
- ▶ Bus transactions (caused by cache controllers)
 - ▶ BusRd: asks for copy with no intent to modify
 - ▶ BusRdX: asks for copy with intent to modify²
 - ▶ Flush: puts data on bus (either requested by another cache or voluntarily due to cache replacement – write back)

²Sometimes BusUpgr is also included to simply ask for permission to modify.

MSI write-invalidate snooping protocol



A / B: if action A is observed by cache controller, action B is taken

---> Broadcast (bus) initiated transaction

—> Processor initiated transaction

Alternative state names:

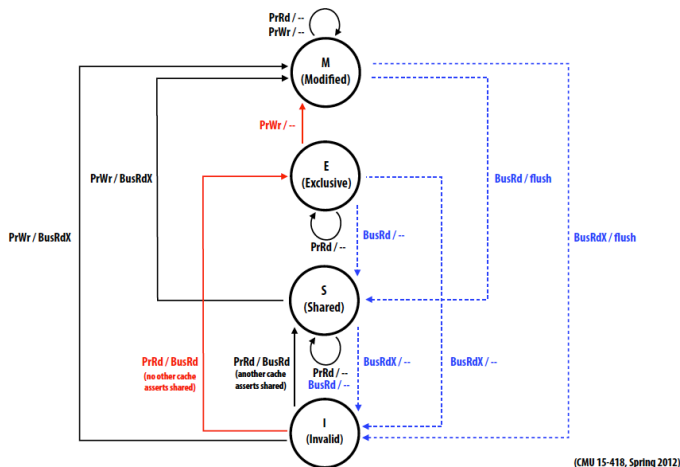
- E (exclusive, read/write access)
- S (potentially shared, read-only access)
- I (invalid, no access)

(CMU 15-418, Spring 2012)

New state for MSI: Exclusive

- ▶ MSI requires two bus transactions for the common case of reading data, and later writing to it
 - ▶ Transaction 1: BusRd to move from I to S state
 - ▶ Transaction 2: BusRdX to move from S to M state
- ▶ This inefficiency exists even if application has no sharing at all
- ▶ Solution: add additional state E (Exclusive clean)
 - ▶ Line not modified, but only this cache has copy
 - ▶ Decouples exclusivity from line ownership (line not dirty, so copy in memory is valid copy of data)
 - ▶ Upgrade from E to M does not require a bus transaction

MESI write-invalidate snooping protocol



MESI: increasing efficiency and complexity

- ▶ Does main memory needs to be updated when flushing?
 - ▶ **MOESI** protocol adds O (Owned, but not exclusive) state: one cache maintains line in O state, other caches maintain shared line in S state
- ▶ Does main memory need to supply data if already in E or S in another cache?
 - ▶ No, but if more than one, which cache should provide it?
 - ▶ **MESIF** protocol adds F (Forward) state: one cache holds shared line in F state rather than S³
- ▶ Cache-to-cache transfers: cache in O or F state is responsible for servicing when required by another cache

³Usually F state migrates to last cache that loads the line, why?

Minimizing sharing

- ▶ True sharing
 - ▶ Frequent writes to a variable can create a bottleneck
 - ▶ Sometimes multiple copies of the value, one per processor, are possible (e.g. the data structure that stores the freelist/heap for malloc/free)
- ▶ False sharing
 - ▶ Cache block may also introduce artefacts: two distinct variables in the same cache block
 - ▶ Technique: allocate data used by each processor contiguously, or at least avoid interleaving in memory
 - ▶ Example problem: an array of ints, one written frequently by each processor (many ints per cache line)

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures


Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Transistors, frequency, power, performance and ... cores!


An inflexion point in 2004 ... the power wall⁴.



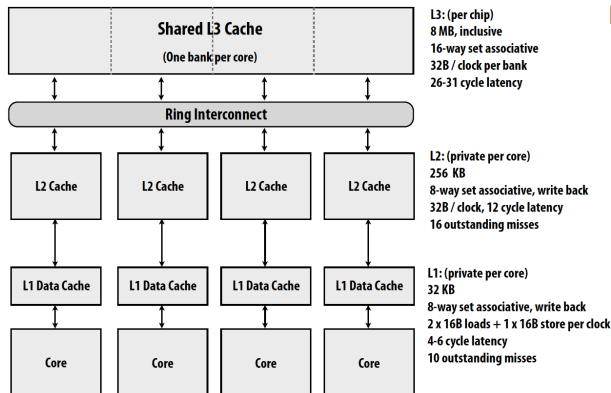
⁴ Original data collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond and C. Batten.

Multicores

- ▶ The increasing number of transistors on a chip is used to accommodate multiple processors (cores) on a single chip
- ▶ Usually private caches (up to a certain cache level) and one last-level cache (LLC)
- ▶ Coherence maintained at the LLC level
- ▶ Chip or socket boundary, access to main memory
- ▶ Multicore = Chip Multi-Processor (CMP)



Example: multicore based on Intel Nehalem i7



Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

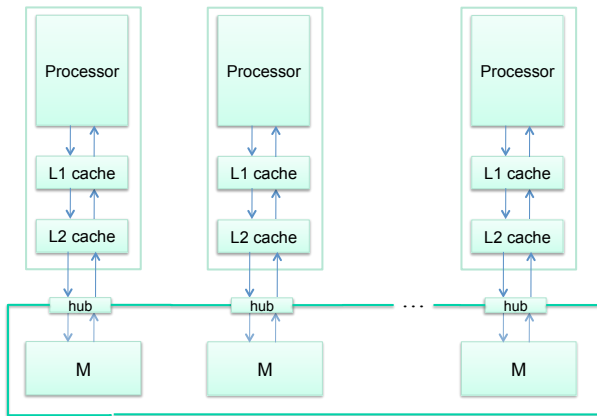
Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Non-Uniform Memory Architectures (NUMA)




- Hub enables cache-coherent NUMA

Directory-based cache coherency

- ▶ Who is involved in maintaining coherence of a memory block?
 - ▶ **Home** node: node (memory of the node) where the block is allocated (OS managed, for example first touch)
 - ▶ **Remote** nodes: **Owner** node containing **dirty** copy or **Reader** nodes containing **clean** copies of the block
 - ▶ **Local** node: node containing the processor requesting the block
- ▶ An additional structure is necessary to track the location of copies of memory block in caches: **Directory**
- ▶ Coherence is maintained by point-to-point messages (not broadcast) between Local/Remote nodes and the directory in the Home node

Directory-based cache coherency

- ▶ Directory structure associated to the node memory: one entry per block of memory
 - ▶ Status bits: they track the state of cache lines in its memory
 - ▶ Sharers list: tracks the list of remote nodes having a copy of a block. For small-scale systems, implemented as a bit string



- ▶ Directory is the centralised structure that "orders" the accesses to each block

Simplified coherency protocol

Possible commands arriving to home node from local node:

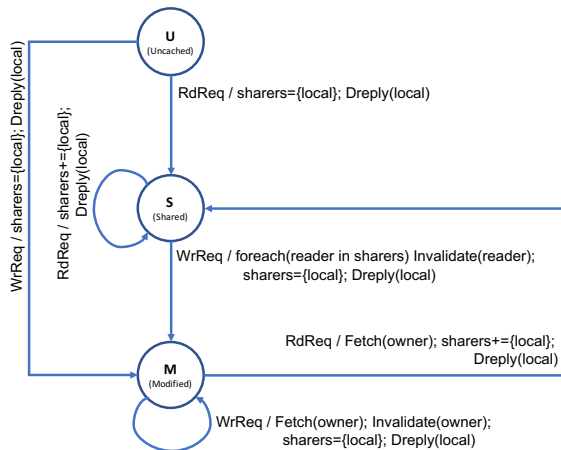
- ▶ RdReq: Asks for copy of block with no intent to modify
- ▶ WrReq: Asks for copy of block with intent to modify (miss in cache of local node) or simply asks for permission to modify it (hit in cache of local node)
- ▶ Dreply: Sends clean copy of block

In response to that, home node⁵ may generate other commands to remote nodes:

- ▶ Fetch: Asks remote (owner) node for a copy of block
- ▶ Invalidate: Asks remote node to invalidate its copy

⁵In other protocols local nodes perform coherence actions based on information provided by the home and owner directly provides data to local.

Simplified coherency protocol



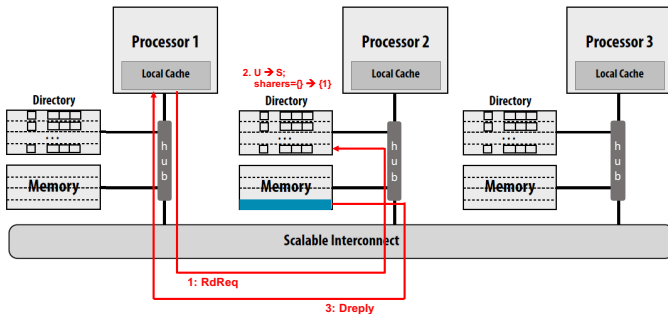
Simplified coherency protocol: additional issues

- ▶ **WriteBack** command: send by local node when cache line is replaced due to cache conflict
 - ▶ Directory status for block transitions from M to U (block in home node needs to be updated), S to S (one less sharer) or S to U (no sharers left after last one)
- ▶ The MSI cache state graph is leveraged to respond to the commands sent from the directory
 - ▶ When **Fetch** is received: line status transitions from M to S and line is flushed
 - ▶ When **Invalidate** is received: line status transitions from S (or M) to I

Directory-based cache coherency: sequence of actions

Example 1: read miss to uncached block

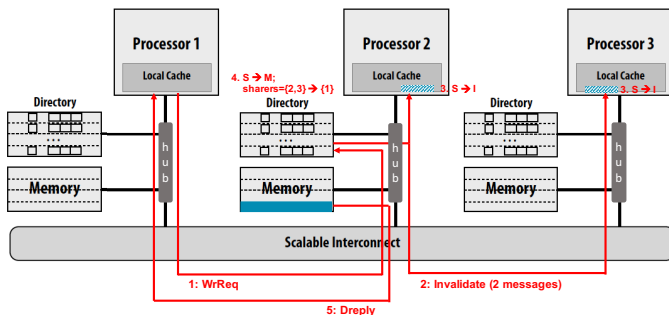
- ▶ Local node where the miss request originates: processor 1
- ▶ Home node where the memory block resides (clean): processor 2



Directory-based cache coherency: sequence of actions

Example 2: write miss to clean block with two sharers


- ▶ Local node where the miss request originates: processor 1
- ▶ Home node where the memory block resides: processor 2
- ▶ Copies of block in caches of processors 2 and 3



Directory-based cache coherency: optimized (optional)

Example 2: write miss to clean block with two sharers


- ▶ Local node where the miss request originates: processor 1
- ▶ Home node where the memory block resides: processor 2
- ▶ Copies of block in caches of processors 2 and 3



Directory-based cache coherency: sequence of actions

Example 3: read miss to dirty block in remote (owner) node

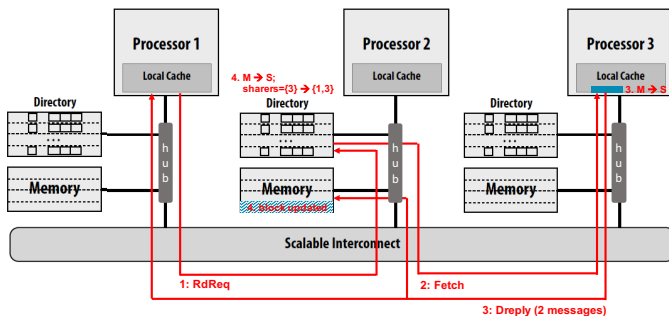
- ▶ Local node where the miss request originates: processor 1
- ▶ Home node for the memory block: processor 2
- ▶ Block dirty currently in cache of processor 3



Directory-based cache coherency: optimized (optional)

Example 3: read miss to dirty block in remote (owner) node

- ▶ Local node where the miss request originates: processor 1
- ▶ Home node for the memory block: processor 2
- ▶ Block dirty currently in cache of processor 3



Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms


The memory consistency problem

Shared memory: address space


Programmer needs

- ▶ Distribute work
- ▶ All threads can access data, heap and stacks
- ▶ Memory is not flat in a NUMA system
 - ▶ True and false sharing even more important
 - ▶ Data allocation and initialization sets the home node
 - ▶ Perform work according to data allocation to minimize data traffic
- ▶ Use synchronization mechanisms to avoid data races

Single process sequential



Single process multithreaded



Why synchronization?

- ▶ Needed to guarantee safety in the access to a shared-memory location or shared resource (e.g. mutual exclusion) or to signal a certain event (e.g. barrier)
- ▶ Components:
 - ▶ Acquire method: how thread attempts to gain access to shared location/resource
 - ▶ Waiting policy: how thread waits for access to be granted to shared location/resource: busy wait, block/awake, wait for a while and then block, ...
 - ▶ Release method: how thread enables other threads to gain access to location/resource once its access completes

Example: a simple, but incorrect, lock

- ▶ What's wrong with ...?
(assume flag initialized to 0, i.e. lock is free; flag equals one means lock is taken)

P1	P2
...	...
lock: ld r1, flag	lock: ld r1, flag
bnez r1, lock	bnez r1, lock
st flag, #1	st flag, #1
... // safe access	... // safe access
unlk: st flag, #0	unlk: st flag, #0
...	...

- ▶ Problem: data race because sequence load–test–store is not atomic!

Support for synchronization at the architecture level

- ▶ Need hardware support to guarantee atomic (indivisible) instruction to fetch and update memory
 - ▶ User-level synchronization operations (e.g. locks, barriers, point-to-point, ...) using these primitives
- ▶ test-and-set: read value in location and set to 1

Example: test-and-set based lock implementation

```
lock:  t&s r2, flag
      bnez r2, lock    // already locked?
      ...
unlock: st flag, #0    // free lock
```

Support for synchronization at the architecture level

- ▶ Atomic exchange: interchange of a value in a register with a value in memory

Example: atomic exchange based lock implementation

```
        daddui r2, r0, #1 // r0 always equals 0
lock:   exch r2, flag      // atomic exchange
        bnez r2, lock      // already locked?
        ...
unlock: st flag, #0       // free lock
```

- ▶ fetch-and-op: read value in location and replace with result after simple arithmetic operation (usually add, increment, sub or decrement)

Support for synchronization at the architecture level

- ▶ Atomicity difficult or inefficient in large systems. Alternative: **Load-linked Store-conditional** ll-sc
 - ▶ ll returns the current value of a memory location
 - ▶ sc stores a new value in that memory location if no updates have occurred to it since the ll; otherwise, the store fails
 - ▶ sc returns success in doing store
- ▶ Examples implementing atomic exchange (left) and fetch-and-increment (right):

```
try: mov r3, r4
      ll r2, location
      sc r3, location
      beqz r3, try
      mov r4, r2
```

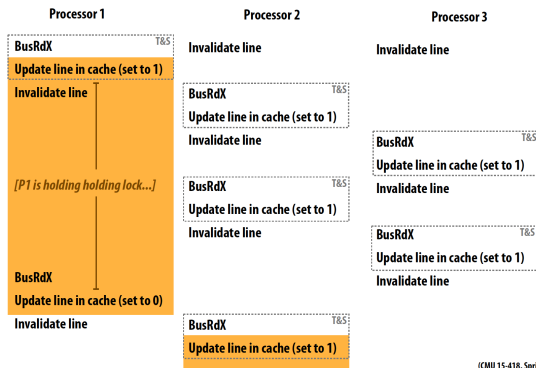
```
try: ll r2, location
      daddui r3, r2, #1
      sc r3, location
      beqz r3, try
```

test-and-set lock coherence traffic

```

lock:  t&s r2, flag      // test and acquire lock if free
      bnez r2, lock     // do it again if already locked
      ...
unlock: st flag, #0     // free the lock

```



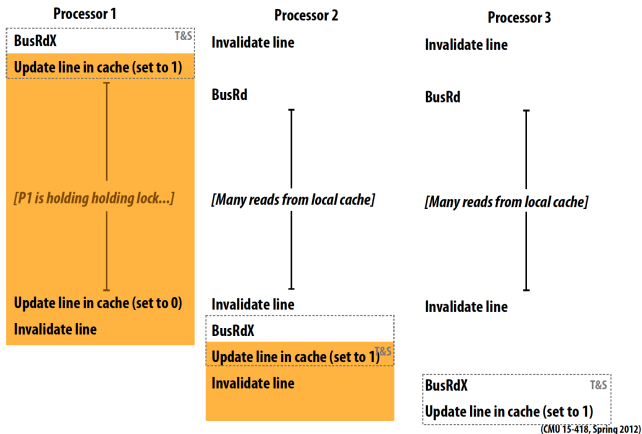
(CMU 15-418, Spring 2012)

Reducing synchronization cost: test-test-and-set

- ▶ test-test-and-set technique reduces the necessary memory bandwidth and coherence protocol operations required by a pure test-and-set based synchronization:
 - ▶ Wait using a regular load instruction (lock will be cached)
 - ▶ When lock is released, try to acquire using test-and-set

```
lock:  ld r2, flag           // test with regular load
      // lock is cached meanwhile it is not updated
      bnez r2, lock          // test if the lock is free
      t&s r2, flag           // test and acquire lock if STILL free
      bnez r2, lock
      ...
unlock: st flag, #0          // free the lock
```


test-test-and-set lock coherence traffic



Reducing synchronization cost: test-test-and-set

- ▶ test-test-and-set technique can also be implemented with ll-sc
 - ▶ First, wait using load linked instruction ll (lock will be cached)
 - ▶ Second, use store conditional sc operation to test if someone else did it first

```
lock:  ll r2, flag           // first test with load linked
      // lock is cached meanwhile it is not updated
      bnez r2, lock         // test if the lock is free
      daddui r2, r0, #1
      sc r2, flag           // try to store 1
      beqz r2, lock         // repeat if someone else did it before me
      ...
unlock: st flag, #0         // free the lock
```

Reducing synchronization cost: test-test-and-set

- ▶ test-test-and-set **idea**⁶ can also help to reduce the synchronization cost of high level parallel programs
 - ▶ Non optimized version : the synchronization is always done

```
acquire_lock(&lock);  
if (value<CONSTANT)    // Test  
    value++;           // Set (Assign)  
release_lock(&lock);
```

- ▶ Optimized version : the synchronization is done if any chance of doing "Set" operation

```
if (value<CONSTANT) {    // Test  
    acquire_lock(&lock); // lock cost is only paid if necessary  
    if (value<CONSTANT)  // Test again  
        value++;         // Set (Assign)  
    release_lock(&lock);  
}
```

⁶Note that `acquire_lock` implementation may also use the test-test-and-set technique to reduce the synchronization cost

Other synchronization primitives

- ▶ How to implement a barrier synchronization primitive?
 - ▶ Threads arriving wait until all have reached the barrier
 - ▶ Structure with fields {lock, counter, flag}

```
barrier:
    acquire_lock(&barr.lock);
    if (barr.counter == 0)
        barr.flag = 0           // reset flag if first
    mycount = barr.counter++;
    release_lock(&barr.lock);

    if (mycount == P) {         // last to arrive?
        barr.counter = 0       // reset counter for next barrier
        barr.flag = 1         // release waiting processors
    } else
        while (barr.flag == 0) // busy wait for release
            ...
```

- ▶ Does it work when consecutive barriers appear? Try to solve it

Outline

Uniprocessor parallelism

Symmetric multi-processor architectures

Multicore architectures

Non-Uniform Memory Architectures

Synchronization mechanisms

The memory consistency problem

Consistency

- ▶ In current systems, the compiler and hardware can freely reorder operations to different memory locations, as long as data/control dependences in sequential execution are guaranteed. This enables:
 - ▶ Compiler optimizations such as register allocation, code motion, loop transformations, ...
 - ▶ Hardware optimizations, such as pipelining, multiple issue, write buffer bypassing and forwarding, and lockup-free caches, ...

all of which lead to overlapping and reordering of memory operations

Consistency: example 1

- ▶ Will writes to different locations be seen in an order that makes sense, according to what is written in the source code?
- ▶ Example: two processors are synchronizing on a variable called flag. Assume A and flag are both initialized to 0

P1	P2
<pre>A=1; flag=1;</pre>	<pre>while (flag==0); /*spin*/ print A;</pre>

- ▶ What value does the programmer expect to be printed?

Consistency: example 2

- ▶ Will writes from one core be seen in a different core, according to what is written in the source code?
- ▶ For example, synchronisation through a shared variable (`next` is implicitly shared):

```
int next = 0;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    for (int end = 0; end == 0; ) {
        ...
        next++;
        if (next==N) end=1;
    }
}
```

```
#pragma omp task
{
    int mynext = 0;
    for (int end = 0; end == 0; ) {
        while (next <= mynext) ;
        ...
        mynext++;
        if (mynext==N) end=1;
    }
}
```


Consistency: example 2 (cont.)

- ▶ Will writes from one core be seen in a different core, according to what is written in the source code?
- ▶ For example, synchronisation through a shared variable (`next` is implicitly shared):

```
int next = 0;
#pragma omp parallel
#pragma omp single
{
    #pragma omp task
    for (int end = 0; end == 0; ) {
        ...
        next++;
        #pragma omp flush(next)
        if (next==N) end=1;
    }
}
```

```
#pragma omp task
{
    int mynext = 0;
    for (int end = 0; end == 0; ) {
        while (next <= mynext) {
            #pragma omp flush(next)
            ;
        }
        ...
        mynext++;
        if (mynext==N) end=1;
    }
}
```

Memory consistency model

The memory consistency model ...

- ▶ Provides a formal specification of how the memory system will appear to the programmer ...
- ▶ ... by placing restrictions on the reordering of shared-memory operations

Sequential consistency, easy to understand but it may disallow many hardware and compiler optimizations that are possible in uniprocessors by enforcing a strict order among shared memory operations.

Memory consistency model

Relaxed consistency (weak), specifying regions of code within which shared-memory operations can be reordered

- ▶ fence machine instruction to force all pending memory operations to complete
- ▶ `#pragma omp flush` and other implicit points in OpenMP language

Different possibilities and implementations to be studied in *Multiprocessors* course

Parallelism (PAR)

Introduction to (shared-memory) parallel architectures

Eduard Ayguadé, Julita Corbalán,
Daniel Jiménez and Gladys Utrera

Computer Architecture Department
Universitat Politècnica de Catalunya

Course 2017/18 (Spring semester)